

The illusion of “agentic AI”: core concepts, limits, and risks made obvious by comparing it with Agency Theory

Riccardo Bovetti

[Who am I?](#) | [LinkedIn](#) | [All you need is thought](#)

February 1, 2026

Abstract

This piece is the latest episode in my recurring series: “*I found a new buzzword and now I have to hoard papers, notes, and hype until my NAS starts judging me.*” The buzzword, today, is *agentic AI*. In the current discourse, “agent” is often a wonderfully convenient label: it sounds serious, sells even better, and makes us believe we are delegating to a digital colleague when, in practice, we are mostly stretching automation—and, in some cases, making it more dangerous.

So I try to take a step back and ask what it really means to *act on behalf of someone*, by contrasting the contemporary enthusiasm for AI agents with an old friend from business economics: *Agency Theory*. If principals have always struggled to control their agents (from the East India companies to modern managers), adding a non-human agent that is opaque and has no intrinsic accountability should not help anyone sleep better.

The goal is not to demonize anything (I do not have the energy), but to clarify concepts, limits, and risks: misalignment, information asymmetries, competence illusions, monitoring costs, and above all that very comfortable temptation to say: “the algorithm decided.” Spoiler: responsibility stays with us.

Contents

1	Operational premise	3
2	Let us start from the beginning: a few historical notions	3
3	A theoretical lens: Agency Theory between principal and agent, and the timeless math of misalignment	4
4	What does “agent” (really) mean—and more specifically, what does it mean in AI? (Or rather: what I think I understood)	6
5	One more step back: introducing the idea of an AI agent in today’s context	7

6	Putting the pieces together: Agency Theory applied to AI agents (and, frankly, to others too)	9
7	Whether we like it or not, AI “agents” are among us and must be managed (after we understand them)	10
8	Emerging risks and latent opportunities	11
9	We are not alone in this dark night (a bibliography-shaped section to stimulate curiosity)	11
10	Conclusions: toward a conscious adoption of agentic AI	12

1 Operational premise

By now it is a recurring pattern: I stumble upon a piece of research, a publication, or an article that promises a new angle, an updated metric, or some further depth on a topic to which, for about two years, I have been dedicating my “research” efforts (i.e., the broad world of AI and its impacts and implications, mostly in business and education). Once I read it, thoughts and images start piling up and force me to chase connections that are sometimes only hinted at.

Over the last year, an enthusiasm (interest?) has spread around the concept of “agentic AI”: AI systems that can act with some form of autonomy, pursuing goals rather than merely executing a process with that goal. A very recent Wharton report sums up this new hype quite bluntly: “agentic AI is the new frontier” and by now *every executive conversation ends up touching on AI agents*.

But given this interest, a provocative question naturally arises: how many of those who use this term have actually reflected on what “agent” means, where the term comes from, and what traps are hidden behind its casual use?

I want to take this space to critically analyze the concept of agentic AI, discussing its theoretical and practical limits (asking, precisely, what “agent” really means in the AI domain—or at least what it meant in November 2025, since recent history teaches us how quickly concepts mutate), and comparing it with the traditional Agency Theory studied and taught in business economics (Jensen & Meckling, Eisenhardt, Ross, etc.).

The underlying idea is to highlight the risks of decision delegation “on behalf”: for the employee and for the company, to an AI system—with organizational, informational, and reputational implications—also considering the role of emerging regulations, with a particular reference to the European AI Act and the broader governance/ethical debates around AI.

2 Let us start from the beginning: a few historical notions

The agents of the Dutch East India Company operating in 17th-century Batavia learned very quickly how to exploit communication delays that could reach up to 18 months with the directors in Amsterdam, creating an information asymmetry that made effective monitoring practically impossible. Agents thus engaged in private trade, appropriating the company’s opportunities, while directors could not verify reports on prices, quality, or market conditions.

The result: weighed down by smuggling, corruption, and rising administrative costs (“compliance” has been with us since time immemorial—or rather, since then), the company collapsed in 1799 after only two centuries.

What, in the last century, came to be called the “agency problem” literally destroyed what had been the largest company in the world.

Similarly, managers of the English East India Company exercised quasi-governmental powers in India far from London’s oversight, taking bribes, manipulating procurement for personal profit, and using the company’s army for private purposes—swimming in diverging time horizons that

encouraged short-term extraction (of value, or rather of valuable goods) rather than long-term value creation.

3 A theoretical lens: Agency Theory between principal and agent, and the timeless math of misalignment

In economic practice (and particularly in management engineering), Agency Theory emerged from observing a pattern repeated across centuries: the problems that arise when a principal delegates the execution of a task or a decision to an agent, in the presence of asymmetric information and potentially divergent interests.

In its modern sense, the theoretical foundation can be traced back to Stephen Ross, who in 1973 formally defined the dynamics of relationships in which “one party acts for, on behalf of, or as a representative of another in a particular domain of decision problems.”

Ross formulated the theory as an embryonic “mathematical challenge” that remains unresolved: when agents maximize their own utility rather than the principal’s welfare, and principals cannot perfectly observe agents’ actions, some residual loss is inevitable—even with optimal monitoring and bonding mechanisms.

The subsequent work of Jensen & Meckling (1976), which defined agency as “a contract under which one or more persons (the principal) engage another person (the agent) to perform some service on their behalf, which involves delegating some decision-making authority to the agent,” made this mathematical foundation sharper.

Agency costs equal monitoring expenditures plus bonding expenditures plus residual loss: the irreducible reduction in welfare stemming from imperfect delegation.

Joseph Stiglitz shared the 2001 Nobel Prize for analyzing markets with asymmetric information, showing that whenever information is imperfect and asymmetrically distributed (which is basically always), economies are not Pareto efficient.

His screening theory showed how uninformed parties design contracts to induce informed parties to reveal private information through self-selection. This became fundamental to understanding insurance markets, credit rationing, and labor relationships.

The challenge intensifies when agents possess both superior information about circumstances and the ability to undertake hidden actions that affect outcomes.

Agency Theory has formalized several categories of problems that arise in these delegation relationships. The main ones (which I cover extensively in my university courses on Management Control, where behavioral components matter to me as much as the technical ones) are:

Adverse Selection a pre-contractual problem, occurring before the agreement is finalized between principal and agent. It is due to ex ante information asymmetry: the agent holds private information about their quality, capabilities, or intentions that the principal does not know, which can lead to a bad choice of agent. In other words, the principal risks selecting, without knowing it, an agent who is less suitable or less honest than they appear, because

the agent hid or distorted information to their advantage. In practice, whoever hires or buys may end up with the “low quality” side of the market if they cannot distinguish quality *ex ante*. In the context of work, an example is hiring a candidate who looks perfect but has skillfully concealed gaps or bad attitudes during the selection process (or boasted fake credentials—which is always best avoided, especially if one is aiming at institutional roles).

Moral Hazard a post-contractual problem, arising after the agency relationship starts. Here, the information asymmetry concerns the agent’s actions once contracted: the agent may behave opportunistically or differently than expected, exploiting the fact that the principal cannot observe or directly control them. In practice, the agent may pursue their own interests at the principal’s expense, especially if they do not bear the consequences. A classic example is the insured party, who after purchasing insurance adopts less prudent behaviors because they know the insurer will cover damages. In the business world, moral hazard is the manager who takes excessively risky decisions because shareholders bear the cost of losses, while the manager benefits (bonuses, stock options) if things go well. In short, the agent is tempted to “shirk” or act non-diligently, knowing the principal struggles to monitor every move.

Misaligned incentives (misalignment) at the root of many agency conflicts lies divergence in objectives or success criteria between principal and agent. If the principal evaluates performance using certain metrics, the agent will optimize their rewards even if that does not coincide with the principal’s maximum benefit. This can produce distortive behaviors. For example, if a salesperson is paid only on commission, they may be tempted to close deals at any cost, because their incentive is to close contracts, while the company would prefer long-term relationships with satisfied customers. Jensen and Murphy (1990) discussed at length incentive alignment, noting that poorly designed compensation systems can worsen agency problems instead of solving them. Agency cost includes precisely the efficiency losses due to suboptimal decisions by the agent because of imperfect incentives, plus the costs of control.

Monitoring and control costs to mitigate the problems above, the principal must invest in mechanisms to monitor the agent’s actions (controls, reporting, audits) and in contracts tying compensation to desired results. All of this generates costs. Shareholders, for example, appoint boards, auditors, internal control committees, invest in information systems, etc. The agent, for their part, may bear bonding costs: expenses to signal commitment or reassure the principal. Even with these measures, there is generally a residual loss: you cannot fully eliminate conflicts without prohibitive costs.

The history of management and moral philosophy reminds us that pursuing an end and following a process that legitimizes that end are not the same thing: there is a deep difference between a rationality of outcomes and a rationality of conduct. The first answers “how do I maximize output?”; the second “which actions are admissible to obtain it?” An ancient

organizational law says: blind optimization of the end produces distortions. In other words, an agent can have instrumental rationality without procedural rationality. A good agent is not only one who reaches the goal, but one who pursues it while respecting the legitimate process—because the process gives shape, transparency, and accountability to the result.

4 What does “agent” (really) mean—and more specifically, what does it mean in AI? (Or rather: what I think I understood)

The term “agentic” applied to AI is often used ambiguously. In the everyday language of the tech industry, an AI agent evokes the idea of an autonomous program that, given an objective, can plan and act independently to achieve it, possibly interacting with the environment and other systems. However, as Ethan Mollick also notes, “agent” remains a fuzzy term: there has been a lot of buzz around agents in recent months, but still (too) little technology that works—or that works really well. In other words: fashionable concept, embryonic implementation.

So we should clarify the foundations of the term before even trying to apply it to AI, so that we do not get dragged by HyperHype in an uncritical way. And yes: when definitions diverge, marketing wins.

The word “agent” carries radically different meanings across disciplines, creating conceptual confusion that marketing tends to exploit systematically. In economics (or more broadly in management), agents are utility maximizers under constraints. Herbert Simon’s work on bounded rationality recognized that real agents “satisfice” rather than optimize due to cognitive limits and information costs, but rationality remains procedural and external. Economic agents do not need to be conscious or autonomous; they simply optimize within constraints.

In philosophy, agency requires intentionality, consciousness, autonomy as self-governance, and moral responsibility. Syntactic manipulation of symbols is insufficient for semantic understanding: algorithms have only “derived” intentionality from their designers, not intrinsic intentionality rooted in consciousness (and it is hard not to mention, here, Searle’s Chinese Room argument).

Traditionally, in classical AI literature (and more properly in the broader Information Technology literature), an agent is any entity capable of perceiving the environment and executing actions upon it. Under this extremely broad definition (Russell and Norvig, 1995—just to remind ourselves we are talking about concepts old enough to vote), even a thermostat is an agent: it “perceives” temperature and “acts” by turning heating on or off.

Obviously, in today’s context, populated by generative AI systems, “agentic AI” is meant as something far more sophisticated. A contemporary perspective emphasizes decision autonomy: one can define an AI agent as a system that uses a large language model to autonomously decide the control flow of an application, i.e., capable of making decisions without step-by-step human instructions.

In practical terms, the difference between an AI tool (or an AI workflow) and an AI agent is the degree of autonomy: a tool executes predefined functions on command; an agent can plan,

reason, and carry out complex tasks independently.

5 One more step back: introducing the idea of an AI agent in today’s context

Let us start with an important distinction: generative AI versus discriminative AI. Discriminative (or analytical) AI is the traditional form focused on data analysis, classification, prediction, or decisions based on rules and predefined models. It is the AI of business intelligence systems and predictive models.

Generative AI, instead, is the family of techniques (from GANs to modern Transformers) that create new content: text, images, audio, code, etc., producing outputs based on patterns learned from training data. In practice: discriminative AI analyzes and explains data; generative AI produces something new.

Often, the most powerful applications combine both approaches: analytical AI extracts information and generative AI uses it to generate useful outputs (like reports or natural-language answers). Each, however, comes with challenges: generative AI can introduce errors or hallucinations if not controlled with robust guardrails; analytical AI is limited to existing patterns and has no “creative” capacity. Understanding this distinction helps clarify what kind of “agents” we are talking about: many so-called AI agents today integrate both analytical and generative components.

Today, when people speak about agentic AI, they typically refer to systems built on top of generative models (LLMs) that are given the ability to interact with external tools (APIs, databases, applications) and to take initiatives to achieve a declared objective. In plain terms, these agents start from a goal, break the task into actions, execute them iteratively, and (sometimes) self-correct in loops.

This paradigm has fired up imaginations: software that autonomously completes complex workflows—from booking travel and meetings, to drafting documents, to managing marketing campaigns—acting as “artificial colleagues”. Often while underestimating two very human organizational constraints that limit agentification: little organizational imagination (limited ability to identify the right use cases) and a lot of organizational mess (undocumented processes, exception-driven work, and heavy reliance on human judgment).

However, reality so far has not met the boldest promises. Many practical implementations of AI agents have shown substantial limits. Often (thankfully), these agents still require strong human intervention (human in the loop) to correct errors, provide feedback, or authorize actions. We are not yet (again: thankfully) facing fully autonomous agents: current systems are closer to extended tools—maybe perceived as more user-friendly thanks to natural language—than to agents with real initiative.

A crucial theoretical (ethical-philosophical) question is: can an AI really have “agency”? In philosophy, agency indicates the capacity to act intentionally and autonomously, making decisions

and taking control of one’s actions. Attributing agency to AI is widely debated. Luciano Floridi observes that today’s AI systems show “a form of agency in specific domains, being able to perform tasks or make decisions within predefined parameters, but they have zero intelligence.” In practice, AI can display goal-oriented agent-like behavior without needing the cognitive properties we normally associate with intelligent action (understanding, consciousness, intentionality).

This “agency without intelligence” comes with practical limits. Generative AI models are famous for the illusion of competence they can create even in the absence of real-world knowledge: they produce plausible, coherent outputs, giving the impression of knowing what they are talking about. In reality, as Emily M. Bender and Timnit Gebru (“Stochastic Parrots”) emphasize, these models are probabilistically recombining language patterns from training data, without understanding meaning. They can output a perfectly formatted, highly plausible, and completely false sentence—without any awareness of error. LLMs do not have an underlying model of the world distinct from the texts they were trained on, nor an intrinsic mechanism to distinguish true from false or to evaluate the appropriateness of actions based on instructions expressed in natural language.

The lack of understanding and common sense makes autonomy fragile. Human action (even if imperfect) has some common sense and causal world-model that helps avoid self-destructive or nonsensical actions while pursuing an objective. Today’s AI agents can get stuck in logical contradictions, be misled by inputs, or attempt inappropriate actions because they cannot truly grasp context. The result is: without supervision, they rarely complete complex tasks reliably.

For years I have been professionally involved in Information Technology, and I have always been critical of the anthropomorphization of technologies (especially digital ones). With AI—discriminative, but especially generative—that process jumped in “quality”: “intelligence”, “learning”, “vision”, “understanding” . . . the list is long.

“Acting” is just the latest human cognitive domain to be colonized terminologically. When mainstream marketing talks about AI agents, it mixes (hopefully intentionally) all those meanings. It promises philosophical agency (autonomy, understanding, intentionality) to justify delegation, while delivering computational agency (automation within constraints).

Anthropomorphization exploits well-known cognitive biases: since Heider and Simmel (1944), we know humans attribute intentionality to anything that shows goal-oriented behavior. Attributing motivations, beliefs, and understanding to AI is what many authors call the “intentional stance”: a useful interpretive strategy, but ontologically empty.

Calling these systems “agents” can lead us to overestimate their capabilities, understanding, and reliability, delegating critical tasks with excessive confidence. Behind apparent agency, AI remains constrained by what it was trained to do. Much of today’s hype is likely a sub-species of AI washing: just as running NLP on a document does not make an application “intelligent”, renaming a bot or a simple RPA script an “agent” does not make it smart.

A true AI agent (which will arrive sooner than later, I am sure) should plan, pursue goals, and collaborate almost like a digital team—but few current solutions get there. Many so-called AI agents today are little more than chatbots or pre-programmed automations slightly boosted

by AI. And in business contexts, full autonomy is often neither possible nor desirable.

In short: handle the notion of agentic AI with care. Yes, AI can act with some autonomy within constraints. No, we do not have artificial intelligences with true understanding, intentionality, or human-level reliability.

6 Putting the pieces together: Agency Theory applied to AI agents (and, frankly, to others too)

With those intrinsic limits in mind, let us examine the idea of considering AI as an “agent” in the economic-organizational sense: a subject acting on behalf of a principal (the company or manager).

Adverse selection in adopting AI when a company chooses to implement an AI system (bought from a vendor or built internally), it may face problems akin to adverse selection. There is ex ante information asymmetry: vendors know their models; the buyer struggles to assess reliability, hidden biases, and technical limitations. Without standard metrics and transparency, the risk is to adopt an inadequate or even counterproductive system.

Moral hazard and AI autonomy it may sound odd to talk about moral hazard with a non-human agent. But if we define moral hazard as the situation in which the agent, once endowed with power to act, takes actions that expose the principal to excessive risk because the agent does not bear the consequences, certain dynamics in AI resemble it. An AI agent does not pay for its own mistakes. If an automated system makes a wrong decision, the company faces economic loss or sanctions; the software keeps running until someone turns it off. This lack of intrinsic accountability can lead AI to pursue the assigned objective without weighing collateral risks (unless tight guardrails are imposed—which may reduce performance).

Objective misalignment (AI alignment) in a traditional principal-agent relationship, the key is structuring incentives and contracts to align the agent with the principal’s goals. In AI, we talk about alignment. How do we ensure the AI optimizes what we really want, rather than a misleading proxy metric? Often, algorithms pursue quantitative objectives (clicks, conversions, measured productivity) that are imperfect surrogates for ultimate goals (service quality, customer satisfaction, reputation). As I always say in my A(I)Wareness courses: be careful what you ask the algorithm—it may satisfy you literally.

Monitoring costs deploying AI agents does not eliminate control costs. If anything, it transforms them, and often initially increases them. We now need skills and structures to monitor AI systems: algorithmic audits, dashboards, validation teams, periodic tests to ensure performance remains within expected bounds. In sensitive applications, we still need a human in the loop: a human operator validating or supervising decisions in real time.

That slows automation and adds personnel costs. But without monitoring, the company may expose itself to severe risks.

Conceptually, then, the use of agentic AI recreates many dynamics typical of the agency relationship: a principal (the organization and its leaders) and an agent (the AI system) with information asymmetry (a black box), with possible opportunistic behaviors (not intentional, but real), and the need to realign and control it.

Not by chance, some scholars explicitly propose analyzing AI implementation in firms through the principal-agent framework. A California Management Review (2025) piece suggests that looking at AI agents as “intelligent workers in service of humans” helps ask the right questions about roles, functions, and governance. That perspective forces us to define oversight mechanisms similar to those used with employees or outsourcers: clear guidelines, performance monitoring, feedback, and continuous correction.

7 Whether we like it or not, AI “agents” are among us and must be managed (after we understand them)

In digitally intensive organizations, an organizational model known as “Hub & Spoke” is spreading. In this model, low value-added, highly standardizable and automatable activities are centralized in a hub, while peripheral units (spokes) specialize in higher-value activities, maintaining strong horizontal connections with the hub to share information and data.

In our context, hub & spoke can be an approach to implementing agentic AI with more control. One could concentrate a “digital brain” in the hub to execute common decisions (e.g., pricing, resource allocation) under close central supervision, while human spoke units in departments interact with the hub, validate outputs, and focus on exceptions or on non-automatable activities.

In short, adopting agentic AI requires careful organizational design to mitigate “algorithmic agency costs”. On one hand, we need principles of system design and monitoring that echo those used to control human agents. On the other, the company may restructure to best exploit AI while retaining control.

In particular, I strongly believe organizations should initially aim for collections of specialized agents collaborating like a team of professionals, each with narrow tasks that are easier to monitor. This facilitates control (each micro-agent is more predictable) while enabling a kind of human-machine collective intelligence.

A practical strategy is to combine generative systems with rule-based control and structured reasoning mechanisms to mitigate erratic behaviors. In other words: “use non-agentic systems to control agentic systems” (as Yoshua Bengio has recently remarked). That is, pair autonomous AI with software guardrails (built in more traditional ways) that verify moves, impose constraints, and block/adjust when deviations occur.

8 Emerging risks and latent opportunities

Critical analysis of business, academic, and societal coexistence with AI in the last three years has made some recurring risk patterns very evident when delegating decisions to algorithms—and these will likely be amplified by the “agentic transformation”.

- **Opacity of motivations:** often, the “why” behind AI decisions is unclear to humans (explainability problem). If the principal does not understand the agent’s actions, the relationship becomes extremely problematic.
- **Human de-responsibilization:** “the algorithm decided” can become an excuse to avoid responsibility. This does not work legally or morally. We must define who is ultimately accountable.
- **Cybersecurity risks:** an AI agent connected to systems can become an attack vector. Delegation to AI introduces new cyber risk surfaces. Inputs must be controlled (prompt injection, etc.) and actions constrained.
- **Impact on people and organization:** delegating decisions affects roles. Some tasks will shrink or change. If automation moves too fast, organizations may lose tacit knowledge.
- **Legal and regulatory aspects:** in some sectors, rules effectively require human involvement in final decisions (thankfully). Delegation boundaries must be traced also based on existing laws and emerging frameworks such as the EU AI Act.

9 We are not alone in this dark night (a bibliography-shaped section to stimulate curiosity)

Many sharp minds are formalizing critical lines of thought that should, at the very least, help bring clarity. Among them, Luciano Floridi stands out for his conceptual distinction between agency and intelligence. In an era of marketing messages about Intelligent Agents, it is refreshing to think of new technology in an anthropocentric key: a useful tool, but one that remains “under control and at the service of an informed society.” Refreshing—and utopian, given how small the informed base is compared to the number of potential users.

Gary Marcus, a long-time skeptic of current AI, has often dismantled the hype around semi-autonomous agents. He argues that current generative architectures lack key properties required for robust autonomous agency, and he frequently highlights hallucinations and non-verifiability.

Emily M. Bender and Timnit Gebru, co-authors of “Stochastic Parrots”, focus on bias, disinformation, and social-linguistic impacts of large language models. Their message is simple: these systems do not understand, and at scale they can amplify harm.

Ethan Mollick, a professor of innovation and entrepreneurship at Wharton, is known for an enthusiastic but reflective adoption of generative AI in work and education. He sees significant

potential, but also points out current limitations: today’s systems are not yet autonomous agents and still need guidance.

10 Conclusions: toward a conscious adoption of agentic AI

Pulling together our path, it becomes evident that “agentic AI” is not a risk-free magic wand, but a powerful tool that requires context, supervision, and wisdom. The initial provocative question—“do we really know what agent means?”—finds an answer: an agent, in the full sense, implies autonomy with responsibility and understanding, which current AI systems do not fully possess.

Today we have semi-autonomous agents that simulate intelligent behaviors in narrow domains, but not agents in the human sense. Treating AI as sentient agents leads to overestimation; treating them as mere software leads to underestimation of impact and unpredictability. We need an enlightened middle way: recognize these systems as hybrid—automatic tools with quasi-agentive capacities—and build around them structures of governance, verification, and human integration.

Agency Theory gives us a useful metaphor: AI is the new agent, and we remain the principals, with the burden of selecting, aligning, monitoring, and taking responsibility for its actions. Where we fail, “agency costs” will show up as wrong decisions, economic losses, reputational damage, or sanctions.

In a sense, training and configuring an AI agent is the new contract drafting: we must codify (in code and data) what we once negotiated verbally with an employee. And just as an incomplete contract can lead to disputes, an algorithm trained on incomplete objectives leads to unexpected outcomes.

Floridi and others remind us that responsibility remains ours: there is no AI to which we can delegate blame. Marcus and others warn us not to be enchanted by the parrot’s song: behind eloquent sentences there may be nobody at home. Bender and Gebru urge us to look at data and real-world impacts to avoid automating injustice. Mollick invites us to experiment—but with our feet on the ground and our hands on the wheel.

For business decision-makers, this translates into a few practical recommendations:

- Learn and understand the technology: do not adopt an AI agent just because everyone does. Invest in AI literacy.
- Start small, evaluate continuously: run limited pilots, measure impact, and scale autonomy based on evidence.
- Define boundaries and fallbacks: establish escalation thresholds, delicate-case blacklists, and kill switches.
- Involve multiple functions: IT, legal/compliance, HR, risk management, and end-user representatives.

- Transparency: communicate what is automated and how; maintain “talk to a human” options.
- Audits and periodic updates: AI is not “set and forget”; it must be monitored and updated.
- Ethics and safeguards: adopt AI ethics guidelines, define accountable roles, and prepare for AI Act compliance.
- Real cost-benefit analysis: do not ignore the cost of monitoring; define clear business cases and measure outcomes.

References

- [1] Wharton Human-Centered AI Report. (2025). *Accountable Acceleration: Gen AI fast-tracks into the enterprise*.
- [2] Jensen, M. C., & Meckling, W. H. (1976). Theory of the firm: Managerial behavior, agency costs and ownership structure. *Journal of Financial Economics*.
- [3] Floridi, L. (2024). Lecture: *Is AI actually intelligent? Has AI agency?* European University Institute.
- [4] Floridi, L. (2025). *La differenza fondamentale*. Mondadori.
- [5] Bender, E. M., Gebru, T., et al. (2021). On the dangers of stochastic parrots: Can language models be too big?
- [6] Mollick, E. (2023). *Almost an Agent: What GPTs can do*. One Useful Thing (blog).